

Buffer overflows

Buffer overflows

- Many CERT security incidents are still buffer overflows
- C and C++ do not perform array bound checks
- Possible to write past the end of an array and
 - Plant malicious code (payload)
 - Hijack control (by overwriting return address)

Buffer overflows

- Call stacks
- Vulnerable functions
- Shellcode
- Stack smashing
- Protection
 - Coding practice
 - Tool support (static, dynamic, and hybrid checks)


Call stacks

- Caller runs

```
push arg1;...; push argN;  
push return_address;
```


- Callee runs

```
push fp;  
fp := sp;  
sp := sp + sizeof(local vars);  
// body of callee  
sp := fp;  
fp := pop();  
pc := pop();
```




Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars);`
`// body of callee`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars) ;`
`// body of callee`
`sp := fp;`
`fp := pop() ;`
`pc := pop() ;`



Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`

- Callee runs

`push fp;`

`fp := sp;`

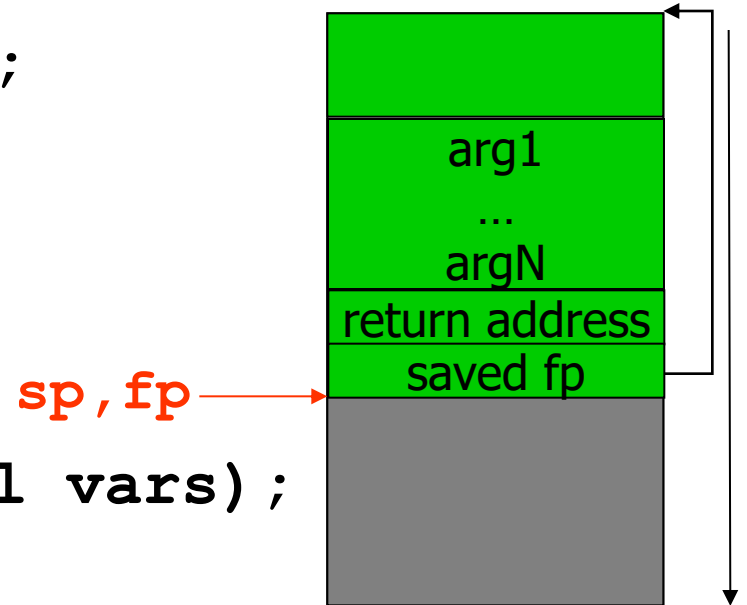
`sp := sp + sizeof(local vars);`

`// body of callee`

`sp := fp;`

`fp := pop();`


`pc := pop();`



Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`


- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars);`
`// body of callee`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Call stacks

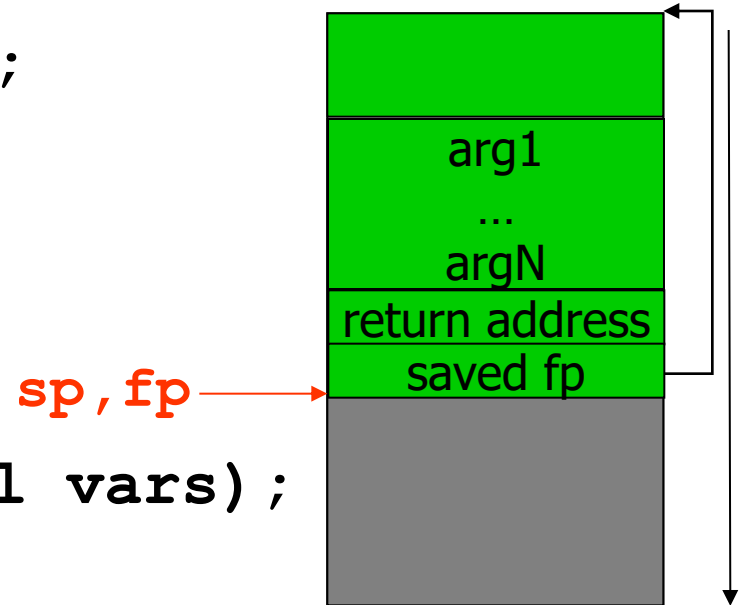
- Caller runs
`push arg1;...; push argN;`
`push return_address;`

- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars) ;`
`// body of callee`
`sp := fp;`
`fp := pop() ;`
`pc := pop() ;`




Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars);`
`// body of callee`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Call stacks

- Caller runs
`push arg1;...; push argN;`
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`sp := sp + sizeof(local vars) ;`
`// body of callee`
`sp := fp;`
`fp := pop() ;`
`pc := pop() ;`



Vulnerable functions

- strcpy, strcat, sprintf, scanf, sscanf, gets, read,...
- No bounds are checked
- Example: gets
 - Reads a buffer from stdin into the buffer
 - No checks for buffer overflows
 - \n (new line) or ^D (EOF) terminate the string

Shellcode

- Shellcode spawns a shell under the uid of the process
 - If uid is elevated to root, this will give rootshell
- How to make sure buffer address overwrites return address?
 - [shellcode][ADDR][ADDR][ADDR]...
- How to make sure return pointer will point to shellcode?
 - No-op sled: [NOP][NOP][NOP]...[shellcode]

Shellcode

- Use gcc and gdb to extract assembly and hex representations
- NOP on the x86 has the machine code 0x90
- How do you guess the ADDR to put in the payload?
 - find out where the first stack frame is: (deterministic in some Linux kernels)
 - offset can be calculated from experiments with different length overruns in gdb
- Attack string example:
 - [NOP][NOP][NOP]...[shellcode][ADDR][ADDR][ADDR]...

Stack smashing

- Vulnerable program:

```
char buf[100];
```


```
...
```

```
gets(buf);
```

- Let us illustrate putting payload instead of the buffer


Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Stack smashing

- Caller runs
`push return_address;`

- Callee runs

`push fp;`

`fp := sp;`

`// allocate space for buffer`

`sp := sp + sizeof(buffer);`

`gets(buffer);`

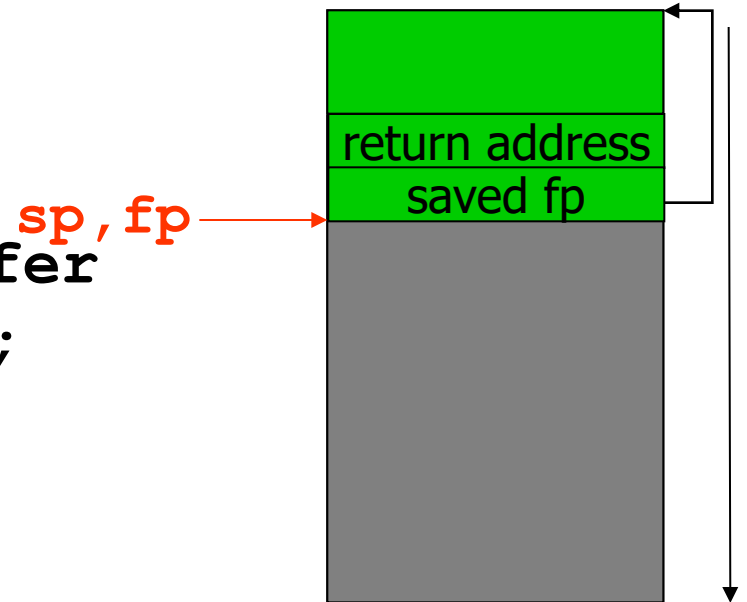
`// user enters shellcode`

`// gets returns`

`sp := fp;`


`fp := pop();`

`pc := pop();`




Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`




Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Stack smashing

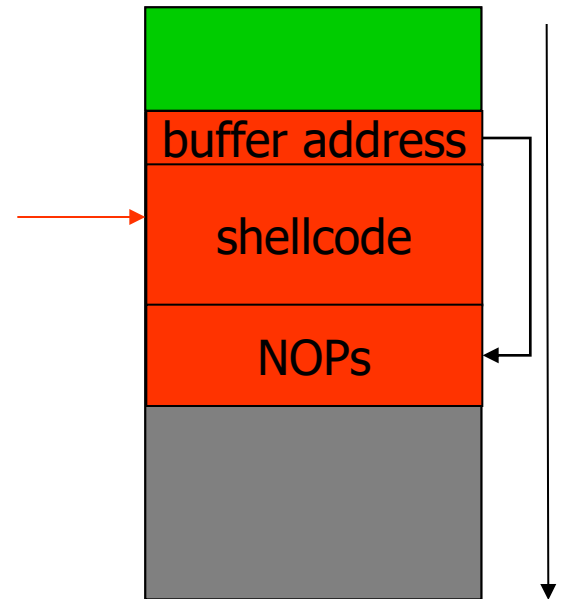
- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Stack smashing

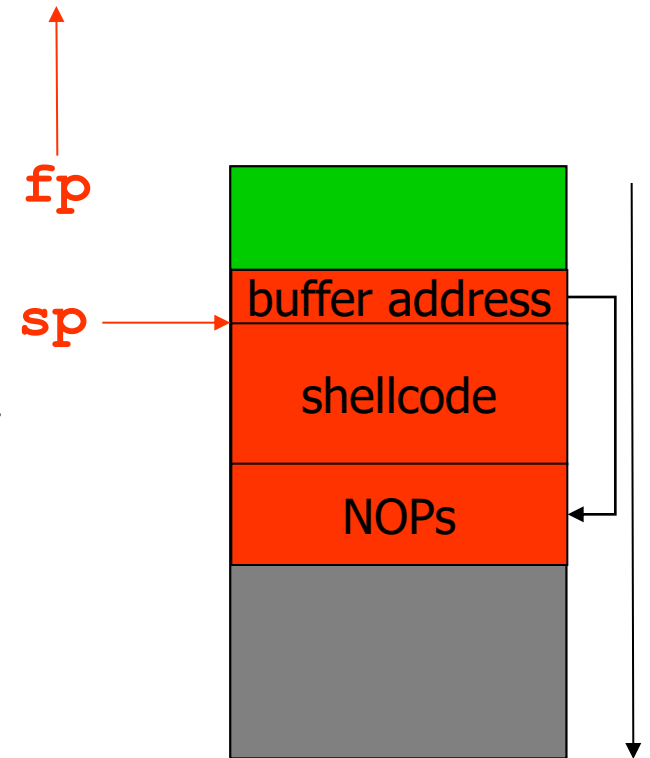
- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`

`sp, fp`



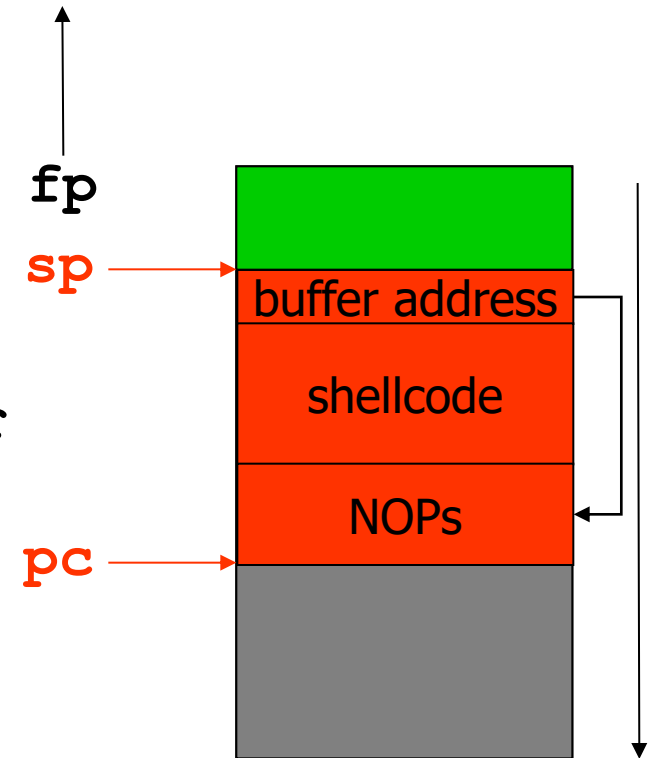
Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Stack smashing

- Caller runs
`push return_address;`
- Callee runs
`push fp;`
`fp := sp;`
`// allocate space for buffer`
`sp := sp + sizeof(buffer);`
`gets(buffer);`
`// user enters shellcode`
`// gets returns`
`sp := fp;`
`fp := pop();`
`pc := pop();`



Protection: coding practice

- Use type-safe languages when you can
- Type-safe dialects of C exist
 - Checked C by Microsoft
 - Need to provide annotations to arrays/pointers
 - Static/dynamic boundary and null pointer checks
 - Rust
 - developed by Mozilla Research
 - static typing
- If you have to use C
 - strcpy -> strncpy,...
 - Restrict the scope of elevated privileges
 - FORTIFY_SOURCE adds boundary checks

Protection: tools

- Static analysis
 - Clang static analyzer, Coverity, ...
- Run-time protection with GCC/Clang
 - -fstack-protector flag adds canaries
 - Clang's SafeStack prevents all stack buffer overflows
- System support
 - Non-executable stack
 - Address obfuscation
 - Memory tagging in hardware

Database security

Database security

- Area by itself
- Basics
- Access control
- Views
- Field protection
- Statistical attacks
- SQL injection
 - see web application security lecture

Basics

- Encrypt connection between caller and db
- User authentication + access control (SQL)
- Protection
 - clearance 0-3

SELECT documents

FROM docdb

WHERE clearance <=1

Access control in SQL

- GRANT action(s) ON object TO user(s)
- REVOKE action(s) ON object TO user(s)
- **Actions:** SELECT, INSERT, DELETE, UPDATE

Views

- “Virtual tables”

```
CREATE VIEW gen_employee_info AS  
  SELECT eid, name, phone, email  
  FROM employee_info
```

- Caution needed

- Views implemented by temporary tables
 - can be slow
- Clients should not connect directly to db
 - o/w a malicious caller may see entire db

Field protection

- Sensitive fields need to be protected besides access controls
 - credit card numbers
- Encryption does the job
 - decrypt on retrieval by authorized party
- How do you search encrypted data?
 - retrieve entire search space and search
- Secure information retrieval
 - direct search on encrypted data
 - search without revealing the query to db

Statistical attacks

- Need to protect individual records
 - salaries, grades,...
- Want to release aggregates
 - average, sum, max, min,...
- How do you prevent statistical attacks?

```
SELECT COUNT(*)  
  FROM students  
 WHERE city = "Åmål"  
    AND course = "concurrent programming"
```

- Suppose result is 1

```
SELECT AVG(grade)  
  FROM students  
 WHERE city = "Åmål"  
    AND course = "concurrent programming"
```

Solution attempt

- Refuse answering queries if they apply to fewer than 10 tuples

```
SELECT COUNT(*)  
  FROM students  
 WHERE course = "concurrent programming"  
::160
```

```
SELECT COUNT(*)  
  FROM students  
 WHERE NOT (city = "Åmål")  
        AND course = "concurrent programming"  
::159
```

- We see that the result is 1

```
SELECT AVG(grade)  
  FROM students  
 WHERE course = "concurrent programming"  
::3.8
```

```
SELECT AVG(grade)  
  FROM students  
 WHERE NOT (city = "Åmål")  
        AND course = "concurrent programming"  
::3.79
```

- $(G_1 + \dots + G_{160}) / 160 = 3.8$ $(G_1 + \dots + G_{160} - G_{\text{Åmål}}) / 159 = 3.79$
- $\therefore G_{\text{Åmål}} = 5$

Defense against statistical attacks

- Analyze statistical inferences
- Restrict query language
 - only average grade for all participants
- Introduce fake tuples
- Give approximate answers
 - Åmål approximated by West Sweden
- Introduce noise
- Differential privacy

TAL: Typed Assembly Languages

TAL: basic type structure

$$type ::= \text{int} \mid \text{code}\{ r_1:t_1, r_2:t_2, r_3:t_3, \dots \}$$

A value with type $\text{code}\{ \mathbf{r1} : t_1, \mathbf{r2} : t_2, \mathbf{r3} : t_3, \dots \}$ must be a label, which when you jump to it, expects you to at least have values of the appropriate types in the corresponding registers.

Simple TAL program with types

```
fact: {r1:int, r2:int, r31:code{r1:int}}  
    ; r1 = n, r2 = accum, r31 = return address  
    sub r3, r1, 1  
    ; {r1:int, r2:int, r31:code{r1:int}, r3:int}  
    ble r3, L2  
    mul r2, r2, r1  
    mov r1, r3  
    jmp fact  
L2: {r2:int, r31:code{r1:int}}  
    mov r1, r2  
    jmp r31
```

Badly typed program

```
fact: {r1:int, r31:code{r1:int}}  
    ; r1 = n, r2 = accum, r31 = return address  
    sub r3, r1, 1; {r1:int, r31:code{r1:int}, r3:int}  
    bge r1, L2  
    mul r2, r2, r1    ; ERROR! r2 doesn't have a  
                      type  
    mov r1, r3  
    jmp fact  
L2: {r2:int, r31:code{r1:int}}  
    mov r1, r2  
    jmp r1    ; ERROR! r1 is an integer!
```

Copyright protection and code obfuscation

Copyright protection

- Tough as any client-side security
- Trade-off with usability and performance
- Impact on privacy (“phone home” applications)
- Impossible to solve but there are tricks to raise the bar for the attacker

Copyright protection schemes

- License keys
 - Can be copied
 - Should not be reconstructable from program (could store key hash)
 - Key checks need to be tamperproof
- On-line Licenses
 - Requires on-line connection (firewall, privacy issues)
 - Checks should be tamperproof
 - Enterprise license servers (against casual pirates)

Tamperproofing

- Antidebugger measures
 - Make debugger crash (via cache)
- Checksums
- Responding to misuse
 - Add subtle bugs
 - Have bugs at start, automatically fix only if no misuse detected
 - Downside: endless support calls...
- Decoys
 - “naïve checks”
 - real checks elsewhere (for example checksums on naïve-check code)
 - Separate copies of license data

Code obfuscation

- Add code that never executes, or does nothing
 - Obvious calculations looking complex
- Move code around
 - Spread related functions
 - Copy and rename the same function instead of calling twice
- Encode your data oddly
 - Strange conversions
 - Encrypt data in memory and hide keys
- Downside
 - Horrible programming style
 - Efficiency
 - Usability/maintanability